

1. Executive Summary

This empirical report presents a comparative evaluation of context routing principles desi

2. Comparative Performance Matrix

Routing Principle Precision Recall F1 Score NDCG@3 Compression % Token S

bm25	100.0%	100.0%	100.0%	0.9000	50.0%	50.0%	80.0%	0.0010	\$0.00
tfidf	100.0%	100.0%	100.0%	0.9000	50.0%	50.0%	80.0%	0.0010	\$0.0
embedding	100.0%	100.0%	100.0%	0.9000	50.0%	50.0%	80.0%	0.0010	

3. Visual Performance Comparisons

![[Retrieval Quality Comparison]](/home/imchaul/Coding Files/Python/maddy_project/data/repor
![[Compression vs Accuracy]](/home/imchaul/Coding Files/Python/maddy_project/data/reports/fi

4. Key Findings & Discussion

- Optimal Router: The bm25 routing principle achieved the highest overall F1 score (100.0%
- Cost-Accuracy Tradeoff: Higher context compression significantly reduces downstream prom
- Latency Impact: Neural and vector similarity routing principles introduce minor initial

5. Conclusion

Context routing principles provide measurable token savings and latency reductions for LLM